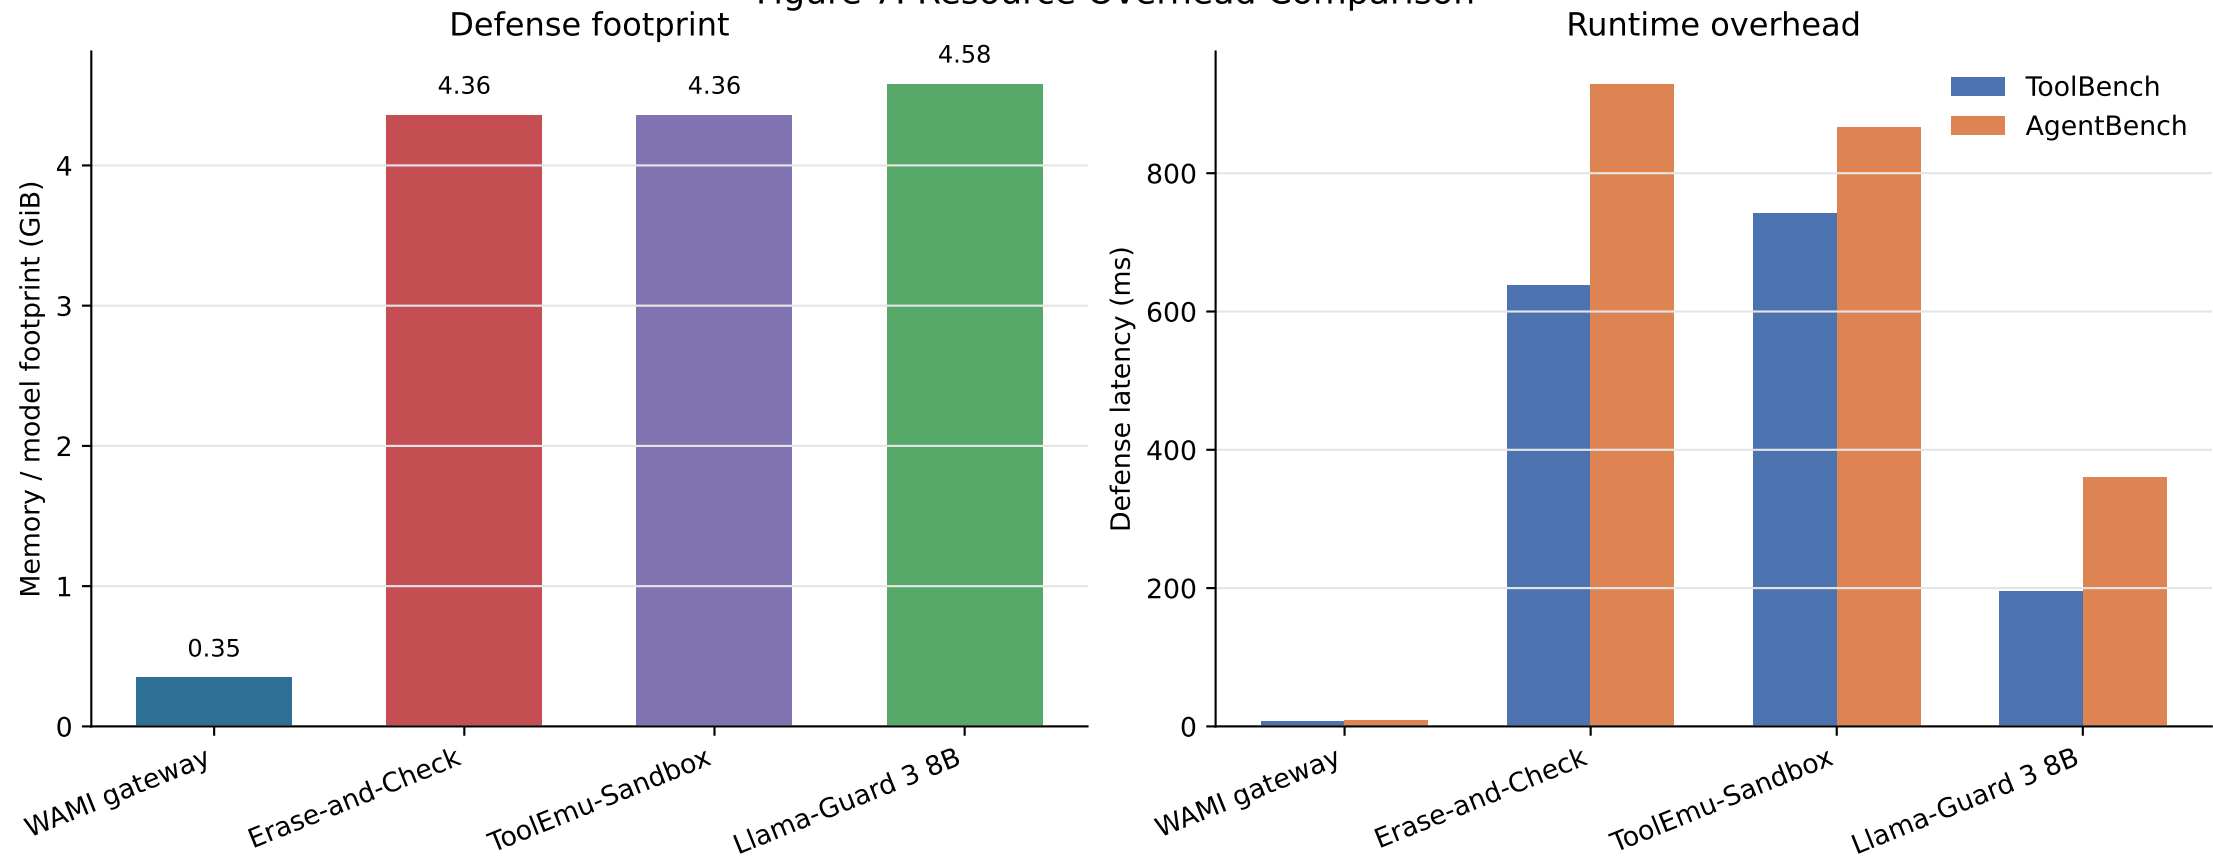


Figure 7. Resource Overhead Comparison



WAMI uses measured CUDA peak reserved memory. LLM baselines use local Ollama model-layer footprint proxy.